

ENGR-E 399/599: Embedded Systems Reverse Engineering

Lecture 3: ELFs and metadata for reverse engineering

Austin Roach
ahroach@iu.edu

January 27, 2022

Mystery function

	0x00000000	0030a0e1	mov r3, r0
→	0x00000004	0320a0e1	mov r2, r3
	0x00000008	0010d2e5	ldrb r1, [r2]
	0x0000000c	013083e2	add r3, r3, 1
	0x00000010	000051e3	cmp r1, 0
<	0x00000014	faffff1a	bne 4
	0x00000018	000042e0	sub r0, r2, r0
	0x0000001c	1eff2fe1	bx lr

- How many arguments does the function take?
- What are the types of the arguments?
- What does the function do?
- What does the function return?

Mystery function revealed

```
size_t strlen(const char *s)
```

Description

The `strlen()` function calculates the length of the string pointed to by `s`, excluding the terminating null byte.

Return value

The `strlen()` function returns the number of bytes in the string pointed to by `s`.

Today's plan

- What is this ELF file format that we've been looking at?
- Why do we care about the format?
- Where did some of this magical metadata (function starts, function names) come from?
- Bonus topics: relocations and dynamic linking

- **Executable and Linkable Format**
 - ▶ Formerly Extensible Linking Format
- File format for executable files, object files, shared libraries, core dumps
- Supports loading program contents into memory and linking

References:

- Tool interface Standard (TIS) Executable and Linking Format (ELF) Specification
- System V Application Binary Interface
- ELF for the ARM architecture
- System V Application Binary Interface AMD64 Architecture Processor Supplement
- Linux Standard Base Core Specification for AMD64
- Additional specifications for other ABIs, minor variants

So why do we care?

It shows up everywhere!

- Linux software (executables and shared objects)
- Embedded Linux systems (executables and shared objects)
- Embedded system boot loaders
 - ▶ Game consoles
- Firmware updates

Why should we study it?

- Source of software reverse engineering metadata
- Non-standard customizations will frustrate your reverse engineering tools
- Other executable formats are likely to have similar elements
 - ▶ Tricks with executable files/loading are plentiful in malware

Other executable file formats

- a.out (Primarily older unixes)
- Mach-O (macOS, iOS)
- PE (Windows, EFI environments)
- Innumerable others

- Most software RE tools are ELF-aware
- Headers parsed in a limited capacity by a large number of tools (such as `file`)
- `readelf`: standard command-line tool for parsing and displaying header information

ELF Header definition

```
#define EI_NIDENT (16)
```

```
typedef struct
```

```
{
```

```
    unsigned char e_ident[EI_NIDENT];    /* Magic number and other info */
    Elf32_Half     e_type;                /* Object file type */
    Elf32_Half     e_machine;             /* Architecture */
    Elf32_Word     e_version;             /* Object file version */
    Elf32_Addr     e_entry;               /* Entry point virtual address */
    Elf32_Off      e_phoff;               /* Program header table file offset */
    Elf32_Off      e_shoff;               /* Section header table file offset */
    Elf32_Word     e_flags;               /* Processor-specific flags */
    Elf32_Half     e_ehsize;              /* ELF header size in bytes */
    Elf32_Half     e_phentsize;           /* Program header table entry size */
    Elf32_Half     e_phnum;               /* Program header table entry count */
    Elf32_Half     e_shentsize;           /* Section header table entry size */
    Elf32_Half     e_shnum;               /* Section header table entry count */
    Elf32_Half     e_shstrndx;            /* Section header string table index */
```

```
} Elf32_Ehdr;
```

readelf -h arm-hello-world

```
Magic:    7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00 00
Class:                                ELF32
Data:                                2's complement, little endian
Version:                            1 (current)
OS/ABI:                              UNIX - System V
ABI Version:                        0
Type:                                DYN (Shared object file)
Machine:                              ARM
Version:                              0x1
Entry point address:                  0x3f4
Start of program headers:              52 (bytes into file)
Start of section headers:              7012 (bytes into file)
Flags:                                0x5000200, Version5 EABI, soft-float ABI
Size of this header:                   52 (bytes)
Size of program headers:               32 (bytes)
Number of program headers:              9
Size of section headers:               40 (bytes)
Number of section headers:              29
```

Program segments

Loader's view of program

- What parts of file should be loaded into memory
- What addresses the segments should be loaded at
- What permissions should be set on the memory segments

Program segment header definition

```
typedef struct
{
    Elf32_Word    p_type;        /* Segment type */
    Elf32_Off     p_offset;      /* Segment file offset */
    Elf32_Addr    p_vaddr;       /* Segment virtual address */
    Elf32_Addr    p_paddr;       /* Segment physical address
                                   (ignored for System V) */
    Elf32_Word    p_filesz;      /* Segment size in file */
    Elf32_Word    p_memsz;       /* Segment size in memory
                                   (>= p_filesz) */
    Elf32_Word    p_flags;       /* Segment flags */
    Elf32_Word    p_align;       /* Segment alignment
                                   (offset % alignment must equal
                                   vaddr % alignment)*/
} Elf32_Phdr;
```


readelf -l arm-hello-world

Program Headers:

Type	Offset	VirtAddr	PhysAddr	FileSiz	MemSiz	Flg	Align
EXIDX	0x00060c	0x0000060c	0x0000060c	0x00008	0x00008	R	0x4
PHDR	0x000034	0x00000034	0x00000034	0x00120	0x00120	R	0x4
INTERP	0x000154	0x00000154	0x00000154	0x00013	0x00013	R	0x1
[Requesting program interpreter: /lib/ld-linux.so.3]							
LOAD	0x000000	0x00000000	0x00000000	0x00618	0x00618	R E	0x10000
LOAD	0x000f08	0x00010f08	0x00010f08	0x0013c	0x00140	RW	0x10000
DYNAMIC	0x000f10	0x00010f10	0x00010f10	0x000f0	0x000f0	RW	0x4
NOTE	0x000168	0x00000168	0x00000168	0x00044	0x00044	R	0x4
GNU_STACK	0x000000	0x00000000	0x00000000	0x00000	0x00000	RW	0x10
GNU_RELRO	0x000f08	0x00010f08	0x00010f08	0x000f8	0x000f8	R	0x1

Memory map! Write/execute permissions of memory segments!

Loading illustrated



Linker's view of program

- Sections of files that can be rearranged/merged during linking
- Information used for dynamic linking

Section header format

```
typedef struct
{
    Elf32_Word    sh_name;        /* Section name (string tbl index) */
    Elf32_Word    sh_type;        /* Section type */
    Elf32_Word    sh_flags;       /* Section flags */
    Elf32_Addr    sh_addr;        /* Section virtual addr at execution */
    Elf32_Off     sh_offset;      /* Section file offset */
    Elf32_Word    sh_size;        /* Section size in bytes */
    Elf32_Word    sh_link;        /* Link to another section */
    Elf32_Word    sh_info;        /* Additional section information */
    Elf32_Word    sh_addralign;   /* Section alignment */
    Elf32_Word    sh_entsize;    /* Entry size if section holds table */
} Elf32_Shdr;
```

readelf -S arm-hello-world

[Nr]	Name	Type	Addr	Off	Size	ES	Flg	Lk	Inf	Al
...										
[5]	.dynsym	DYNSYM	000001c4	0001c4	0000a0	10	A	6	3	4
[6]	.dynstr	STRTAB	00000264	000264	000086	00	A	0	0	1
...										
[12]	.plt	PROGBITS	000003a4	0003a4	000050	04	AX	0	0	4
[13]	.text	PROGBITS	000003f4	0003f4	0001fc	00	AX	0	0	4
[14]	.fini	PROGBITS	000005f0	0005f0	000008	00	AX	0	0	4
[15]	.rodata	PROGBITS	000005f8	0005f8	000011	00	A	0	0	4
...										
[21]	.got	PROGBITS	00011000	001000	00003c	04	WA	0	0	4
[22]	.data	PROGBITS	0001103c	00103c	000008	00	WA	0	0	4
[23]	.bss	NOBITS	00011044	001044	000004	00	WA	0	0	1
...										
[26]	.symtab	SYMTAB	00000000	001098	000690	10		27	81	4
[27]	.strtab	STRTAB	00000000	001728	000337	00		0	0	1
...										

Some common sections

.text	Executable instructions of a program
.data	Initialized, writable data
.rodata	Initialized, read-only data
.bss	Uninitialized data
.symtab	Symbols table
.strtab	Strings table (typically for .symtab entries)
.dynsym	Dynamic linking symbols table
.dynstr	Strings for dynamic linking
.got	Global offset table
.plt	Procedure linkage table

We know where executable code is! Where data is!

Section to segment mapping

Section to Segment mapping:

Segment Sections...

00	.ARM.exidx
01	
02	.interp
03	.interp .note.gnu.build-id .note.ABI-tag .gnu.hash .dynsym .dynstr .gnu.version .gnu.version_r .rel.dyn .rel.plt .init .plt .text .fini .rodata .ARM.exidx .eh_frame
04	.init_array .fini_array .dynamic .got .data .bss
05	.dynamic
06	.note.gnu.build-id .note.ABI-tag
07	
08	.init_array .fini_array .dynamic

Symbols

Support linking:

- Symbols used to refer to functions, variables, etc. by name
- Linker resolves symbols to match addresses to resources

Regular symbols

Used at build time, and can be discarded from the final executable or shared library

Dynamic symbols

Used for runtime dynamic linking, and cannot be discarded without breaking this process

Symbol table entries

```
typedef struct
{
    Elf32_Word    st_name;        /* Symbol name (string tbl index) */
    Elf32_Addr    st_value;      /* Symbol value */
    Elf32_Word    st_size;       /* Symbol size */
    unsigned char st_info;       /* Symbol type and binding */
    unsigned char st_other;      /* Symbol visibility */
    Elf32_Section st_shndx;      /* Section index */
} Elf32_Sym;
```

Symbol examples

```
$ readelf -s arm-hello-world
```

```
Symbol table '.dynsym' contains 10 entries:
```

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
...							
5:	00000000	0	FUNC	GLOBAL	DEFAULT	UND	puts@GLIBC_2.4 (2)
...							

```
Symbol table '.symtab' contains 105 entries:
```

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
...							
101:	00000570	28	FUNC	GLOBAL	DEFAULT	13	main
...							

Names, addresses, and lengths of functions and global variables!

Dynamic symbols

Support *dynamic linking* of shared-object libraries on an as-needed basis.

- Needed libraries identified in the 'dynamic' section (`readelf -d`)
- Dynamic symbols used to identify functions or variables that will be requested by a program or library
- Dynamic symbols also used to identify functions or variables that are made available by a library
- Dynamic linker identifies requested functions or variables by name

Stripped binaries

- Symbol table entries are extraneous in fully linked executables
- Symbols can take up a lot of space
- Symbols can be stripped with `strip` command
- Dynamic symbols can *not* be stripped without affecting the functionality of the executable
 - ▶ Dynamic symbols are used for runtime dynamic linking of shared object files

Consequences of stripping

	Unstripped	Stripped
Dynamically linked	All symbols and dynamic symbols intact	Only dynamic symbols
Statically linked	All symbols intact (no dynamic symbols)	No symbols

Finding main() in a stripped Linux ELF

- Entry address is *not* the address of main()
- Typical to call __libc_start_main() as first function call

```
int __libc_start_main
(int *(main) (int, char **, char **), // Address of main()
 int argc,                          // argc
 char * * ubp_av,                    // char *argv
 void (*init) (void),                // Initialization routines
 void (*fini) (void),                // Termination routines
 void (*rtld_fini) (void),           // Dynamic linker's finalizer
 void (* stack_end));                // Stack pointer (?)
```

Application examples: stripped versus unstripped

For reference, `arch/arm/tools/syscall.tbl` in the Linux kernel source tree:

0	common	restart_syscall	sys_restart_syscall
1	common	exit	sys_exit
2	common	fork	sys_fork
3	common	read	sys_read
4	common	write	sys_write
5	common	open	sys_open
6	common	close	sys_close
...			

Finding function entry points in absence of symbols

- Defined entry points (reset/interrupt vectors)
- Targets of call instructions
- Function preamble heuristics

Bonus material

Debugging symbols

- DWARF standard format for Linux ELF's
- Encoded in various `.debug` sections
- Provides information about variable types, mapping of instructions to source code lines, etc.
- Typically quite large. Very uncommon to find these in production software.

You will probably not be this lucky

Why do these addresses keep changing?

```
int foo = 1;
int main(int argc, char **argv) {
    printf("Address of foo: %p\n", &foo);
    return 0;
}
```

```
$ ./print-pointer
```

```
Address of foo: 0x4ce044
```

```
$ ./print-pointer
```

```
Address of foo: 0x42a044
```

```
$ ./print-pointer
```

```
Address of foo: 0x475044
```


Relocations

To support relocating code (ASLR, dynamically loaded shared libraries), values in memory (especially pointers) can be specified as relocatable and will be modified at load time.

This is primarily a concern if you want to patch the binary.

Relocations

```
typedef struct
{
    Elf32_Addr  r_offset;    /* Address */
    Elf32_Word  r_info;      /* Relocation type and symbol index */
} Elf32_Rel;
```

```
typedef struct
{
    Elf32_Addr  r_offset;    /* Address */
    Elf32_Word  r_info;      /* Relocation type and symbol index */
    Elf32_Sword r_addend;    /* Addend */
} Elf32_Rela;
```

Dynamic linking mechanism

Relevant sections:

- `.plt`: Procedure linkage table
 - ▶ Stubs to look up and jump to entries in GOT
- `.got`: Global offset table
 - ▶ Addresses of external symbols filled in dynamically by linker
 - ▶ Sometimes a special GOT just for the PLT (`.got.plt`)

This is primarily a concern for the development of software exploits.

Example

```
(gdb) b main
(gdb) run
(gdb) x/3i $pc
=> 0x400518 <main+12>: add r3, pc
    0x40051a <main+14>: mov r0, r3
    0x40051c <main+16>: blx 0x4003cc <puts@plt>
```

The program wants to call a dynamically linked function, `puts()`. It calls the associated PLT stub.

Example

```
(gdb) x/3i 0x4003cc
0x4003cc <puts@plt>: add r12, pc, #0, 12
0x4003d0 <puts@plt+4>: add r12, r12, #16, 20 ; 0x10000
0x4003d4 <puts@plt+8>: ldr pc, [r12, #3132]! ; 0xc3c
(gdb) x/1xw 0x411010
0x411010 <puts@got.plt>: 0x004003ac
```

The called PLT stub fetches an address from the GOT and jumps to it.

Example

```
gdb) x/5i 0x4003ac
0x4003ac: push {lr} ; (str lr, [sp, #-4]!)
0x4003b0: ldr lr, [pc, #4] ; 0x4003bc
0x4003b4: add lr, pc, lr
0x4003b8: ldr pc, [lr, #8]!
(gdb) x/1xw 0x4003bc
0x4003bc: 0x00010c44
```

The first time, the PLT stub invokes the dynamic linker. The GOT entry address is still stored in r12, which indicates the desired function to the linker.

Example

```
(gdb) b *0x400520
```

```
Breakpoint 2 at 0x400520
```

```
(gdb) c
```

```
Continuing.
```

```
Hello world!
```

```
Breakpoint 2, 0x00400520 in main ()
```

The dynamic linker executes `puts()` and fixes up the GOT entry behind the scenes.

Example

```
(gdb) x/1xw 0x411010
0x411010 <puts@got.plt>: 0xb6f24ab1
(gdb) x/5i 0xb6f24ab1
0xb6f24ab1 <puts>: stmdb sp!, {r4, r5, r6, r7, r8, lr}
0xb6f24ab5 <puts+4>: mov r6, r0
0xb6f24ab7 <puts+6>: bl 0xb6f35380 <strlen>
0xb6f24abb <puts+10>: ldr r5, [pc, #316] ; (0xb6f24bf8 <puts+328>)
0xb6f24abd <puts+12>: ldr r2, [pc, #316] ; (0xb6f24bfc <puts+332>)
```

And now the puts() GOT entry points directly to puts() in libc!

Dynamic linking illustration

