

ENGR-E 399/599: Embedded systems reverse engineering

Lecture 8: Microcontroller emulation

Austin Roach
ahroach@iu.edu

March 3, 2022

Mystery function

	0x00000000	ff1001e2	and r1, r1, 0xff
>	0x00000004	0020d0e5	ldrb r2, [r0]
	0x00000008	0030a0e1	mov r3, r0
	0x0000000c	010052e1	cmp r2, r1
	0x00000010	010080e2	add r0, r0, 1
<	0x00000014	0200000a	beq 0x24
	0x00000018	000052e3	cmp r2, 0
<	0x0000001c	f8ffff1a	bne 4
	0x00000020	0230a0e1	mov r3, r2
>	0x00000024	0300a0e1	mov r0, r3
	0x00000028	1eff2fe1	bx lr

- How many arguments does the function take?
- What are the types of the arguments?
- What does the function do?
- What does the function return?

Mystery function revealed

```
char *strchr(const char *s, int c)
```

Description

The `strchr()` function returns a pointer to the first occurrence of the character `c` in the string `s`.

Return value


The `strchr()` function returns a pointer to the matched character or `NULL` if the character is not found. The terminating null byte is considered part of the string, so that if `c` is specified as `'\0'`, this function returns a pointer to the terminator.

Today's plan

- Discussion of brute-forcing approaches
 - ▶ Speeding up UART brute-forcing with firmware patching
- Emulation overview
- Introduction to `simavr`
 - ▶ Very relevant to Assignment 03
- Emulation for dynamic analysis
- Popular emulation frameworks

Reminder: brute-forcing hash functions


Assignment 3 contains a cryptographic hash function:



Goal is to determine an input that hashes to a specific value.

To find a match for a good cryptographic hash function, the best you can do is test many possible inputs (brute force).

Brute-forcing bottleneck




Most of our time is spent sending characters back and forth across the UART

Speeding up brute-forcing over UART

- Speed up UART interface
 - ▶ 38400 to 57600 baud; faster would require some platform modifications
- Modify strings to print fewer characters
 - ▶ Truncate each message to a single character and newlines
- Start sending next password before result is received
- Run generator and evaluation routine on the microcontroller

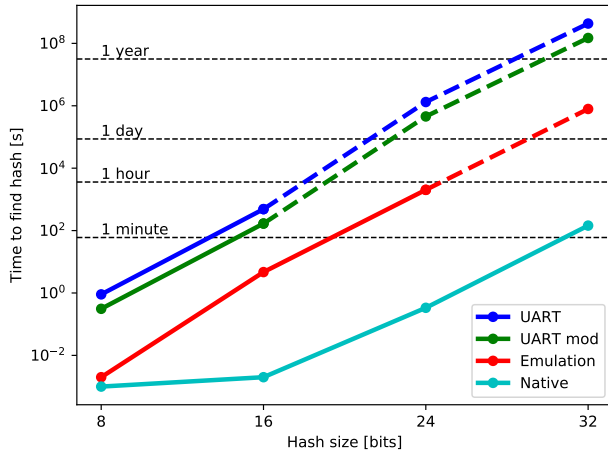
Let's patch some firmware!

Other approaches



- Emulate the code on a faster system
 - ▶ Today's topic
- Analyze the hash algorithm, re-write and compile to run on another, faster processor
 - ▶ You'll get to do this in Assignment 3

Other approaches compared



Emulation overview

- Emulation allows one system to execute code meant for another system
 - ▶ Often cross-architecture
- We have previously emulated using QEMU
 - ▶ Linux 'user-mode' emulation of the binary water balloon
 - ▶ System-level emulation of embedded Linux systems
- Can also use emulators to execute sequences of instructions
 - ▶ Not emulating the full system
 - ▶ Carve out the pieces of the code that you care about

Emulation applications

- Dynamic analysis when we don't have debug access on the host system
- Incorporate pieces of cross-architecture code in larger execution framework
- Parallelize searching tasks
- Virtualize I/O or peripheral interaction to investigate system behaviour under various conditions

Emulation frameworks

- Keep track of processor/system state
- Model the effects of instructions on the system state
- Peripheral interaction:
 - ▶ Simulate
 - ▶ Pass-through to real hardware
 - ▶ Provide callbacks for user implementation

Modeling peripheral interaction typically the biggest challenge for embedded systems emulation

Processor state in simavr

Represented by `avr_t` struct

- https://github.com/buserror/simavr/blob/master/simavr/sim/sim_avr.h
- Configuration settings (clock frequency, voltages, fuse settings)
- Device information (memory sizes, address size)
- Memory (including register states)
- SREG state
- Handler functions for special events
- Interrupt handler state
- Statistics
 - ▶ Cycle counts
- Emulation state (cycle limits, termination address)

Instruction modeling in simavr

Implemented in `avr_run_one()`:

- https://github.com/buserror/simavr/blob/master/simavr/sim/sim_core.c
- Parse machine code representation
 - ▶ Determine opcode
 - ▶ Determine operands
- Advance program counter
 - ▶ Add 2 by default
 - ▶ Control flow instructions alter in other ways
- Update memory or register file
- Update status register

simavr: initialization

```
avr_t *avr;  
avr = avr_make_mcu_by_name("atmega328p");  
avr_init(avr);  
// Define frequency, voltages, etc.
```

simavr: program memory

```
typedef struct avr_t {  
    ...  
    uint8_t * flash;  
    ...  
}
```

Initialized in our example with:

```
memcpy(avr->flash, prog_mem, _PROG_MEM_SIZE);
```


simavr: data memory

```
typedef struct avr_t {  
    ...  
    uint8_t * data;  
    ...  
}
```

Can be used to access general-purpose registers, I/O control registers, and SRAM.

```
// Set register R0 contents to 0x12  
avr->data[0] = 0x12;  
// Set memory address 0x200 to 0x30  
avr->data[0x200] = 0x30;
```

simavr: control

```
// Define a starting address
avr->pc = STARTING_ADDRESS;
// Run until ending address
int state = avr->state;
while ((state != cpu_Done) && (state != cpu_Crashed) &&
      (avr->pc != ENDING_ADDRESS)) {
    state = avr_run(avr);
}
```

simavr treats the PC as a *byte* address

simavr: skeleton framework for assignment

simavr: Debug

- Can optionally create a GDB debugger stub
- Emulator will wait for debugger to attach before execution

```
avr->gdb_port = 1234;  
avr_gdb_init(avr);
```

Ubuntu packages

```
libsimavr2, libsimavr-dev, gdb-avr
```

simavr: Executing code with debug

```
$ ./emulate uart_intro.bin  
Read 380 bytes from uart_intro.bin.  
avr_gdb_init listening on port 1234  
gdb_network_handler connection opened
```

simavr: Connecting with GDB

```
$ avr-gdb
```

```
...
```

```
(gdb) target remote localhost:1234
```

```
Remote debugging using localhost:1234
```

```
warning: No executable has been specified and target does not support  
determining executable automatically. Try using the "file" command.
```

```
0x00000000 in ?? ()
```

simavr: Setting breakpoints in GDB

```
(gdb) b *0x13a
```

```
Breakpoint 1 at 0x13a
```

```
(gdb) c
```

```
Continuing.
```

Note: automatically using hardware breakpoints for read-only addresses.

```
Breakpoint 1, 0x0000013a in ?? ()
```

```
(gdb)
```

- Code addresses are *byte* addresses in this context
- Some versions of avr-gdb require a crazy cast to set breakpoints at program memory addresses:

```
b *(void (*)())0x13a
```

simavr: Inspecting system state in GDB

```
(gdb) info reg
r0          0x0          0
r1          0x0          0
r2          0x0          0
r3          0x0          0
...
(gdb) x/2s 0x800100
0x800100: "Please enter the passkey:\r\n"
0x80011c: "SUCCESS!\r\n"
...
(gdb) x/1xb 0x800024
0x800024: 0x20
```

GDB adds 0x800000 to SRAM addresses to separate Harvard architecture address spaces.

simavr: Examining instruction effects in GDB

```
(gdb) info reg
```

```
...
```

```
r24          0x18          24
```

```
r25          0x0           0
```

```
...
```

```
(gdb) si
```

```
0x0000013c in ?? ()
```

```
(gdb) info reg
```

```
...
```

```
r24          0x0           0
```

```
r25          0x0           0
```

```
...
```

```
(gdb) si
```

```
0x0000013e in ?? ()
```

```
(gdb) info reg
```

simavr: Dumping memory state with GDB

```
(gdb) dump memory /tmp/mem.bin 0x800000 0x800900
```

```
$ xxd mem.bin
```

```
00000000: 0000 0000 0000 0000 0000 0000 0000 0000  ....
00000010: 0001 0000 0000 0000 0001 2801 ff08 7c01  ....(...|.
00000020: 0000 0000 2000 0000 0000 0000 0000 0000  ....
00000030: 0000 0000 0000 0000 0000 0000 0000 0000  ....
00000040: 0000 0000 0000 0000 0000 0000 0000 0000  ....
...
```

simavr: Importing memory into Ghidra

- Delete SRAM block in Memory Map
- Use File→Add to Program...
- Map 0x800 bytes from offset 0x100 in file to mem:0x100

Quick EMUlator <https://www.qemu.org/>

- User-mode emulation
- System-mode emulation
- Supports (at least): Aarch64, Alpha, ARM, AVR, CRIS, HPPA, i386, m68k, Microblaze, MIPS, MIPS64, OpenRISC, PowerPC, RISC-V, s390x, SH4, SPARC, SPARC64, Tricore, unicore32, x86-64, xtensa, TILE-Gx

Unicorn engine

<https://www.unicorn-engine.org/>

- Uses QEMU for instruction emulation
- Emulate raw binary code (not an ELF or disk image)
- Bindings for multiple programming languages
- Register all sorts of customized handlers
- Falling behind QEMU a bit...

Translation

Emulation is still not the fastest way to calculate these hashes:

- 32-bit arithmetic with 8-bit instructions
- Instruction translation slower than native instruction execution

Translation to your PC's native ISA is significantly faster:

- Understand the hash function algorithm
- Write it in C (or another language)
- Compile for the processor on your PC's processor (probably x86)
- Profit!

To do this successfully requires an exact understanding of the hash algorithm